

Distributed programming with Java

Distributed programming with Java

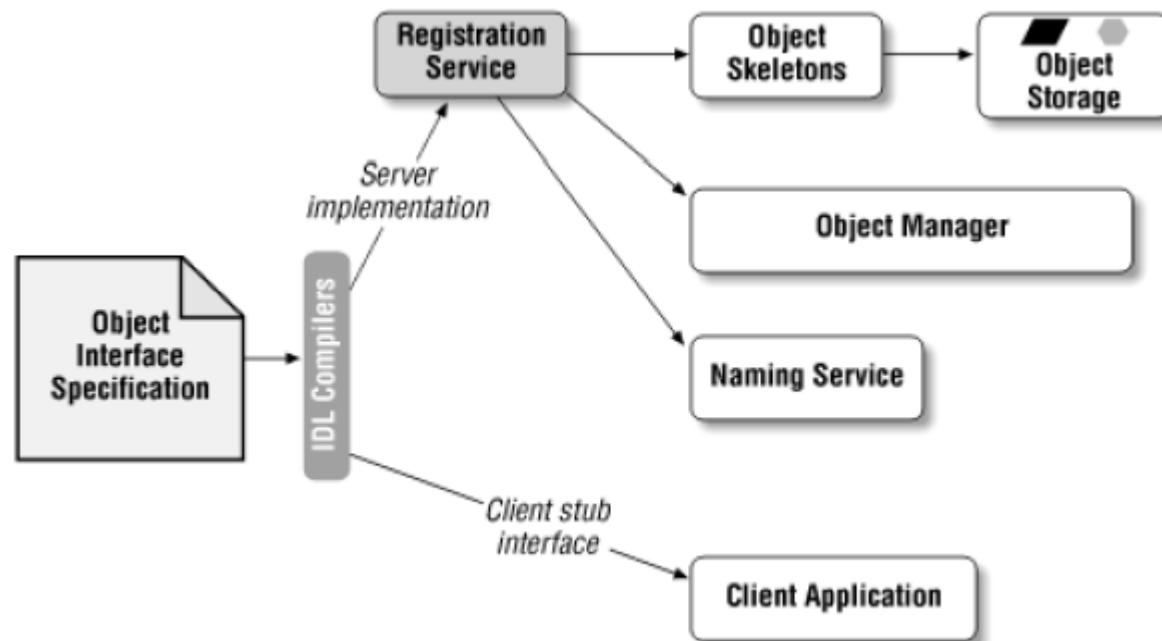
- Introduction
- RMI
- Serialization
- The Reflection API
- Case study – mobile agents

Introduction

- *Distributed programming* is the process of breaking down an application into individual computing objects that can be scattered across a network of computers, yet still work together to do cooperative tasks
- Main advantages of distributed computing include:
 - *Parallelism*: using smaller, cheaper computers to solve large problems instead of resorting to large computers
 - *Reduced network bandwidth*: large data sets are typically difficult to relocate, and easier to control and administer where they are, so remote data servers can be used to provide needed information
 - *Redundancy*: computing objects on multiple networked computers can be used by systems that need fault-tolerance; if a machine goes down, the job can still carry on

Distributed object systems

- *Distributed object system* (DOS) represents a set of APIs that hides the complexity of distributed programming
- It allows programmers to invoke objects on remote hosts, and interact with them as if they were objects within the local host
- A general architecture of a distributed object system:



DOS features

- *Object interface specification*: provides the means (e.g. a language) for specifying object interfaces, regardless of the implementation details
- *Stubs*: responsible for dispatching remote requests, stubs are used to route local method invocations to the object on a server
- *Skeletons*: responsible for processing remote requests, skeletons are used by servers to create new instances of remote objects and to route remote method calls to the object implementation
- *Naming service*: associates a remote object with a name that clients can use to obtain a reference to the object

DOS features (cont')

- *Object manager*: the heart of the distributed object system, it manages the object skeletons and object references on an object server
 - Usually, it also supports more advanced features, such as object persistence
- *Registration service*: registers newly implemented classes with a naming service and an object manager, and then stores the class in the server's storage
- *Object communication protocol*: supports the means for transmitting and receiving object references, method references, and data
 - Ideally, the client application does not know any details about this protocol

Distributed programming with Java

- Introduction
- RMI
- Serialization
- The Reflection API
- Case study – mobile agents

RMI

- *Remote Method Invocation* (RMI) is a core Java API and class library that allows Java programs running in one JVM to call methods in objects running in a different, possibly physically distributed, JVM
 - RMI creates the illusion that this distributed program is running on one system with one memory space
- It is the pure Java approach to distributed software development
- For its functionality, RMI relies on advanced Java features, such as *object serialization*, *reflection*, and *dynamic class loading*
- The RMI API is available under the *java.rmi* package

RMI services

- A *RMI service* is a remote object with methods that may be invoked from a different JVM than the one in which the object itself lives
- Each RMI service implements the *remote interface* that specifies which of its methods can be invoked by clients
- From the programmer's perspective, RMI services work pretty much like the local objects
 - That is, clients invoke the methods of the remote object almost exactly as they invoke local methods
- However, RMI is much slower and less reliable than regular local method invocation
 - Things can and do go wrong with RMI that do not affect local method invocations

The RMI registry

- In order to allow clients to find them, RMI services must be registered with a lookup service called *RMI registry*
- The registry runs as a separate process and allows applications to register RMI services or obtain a reference to a named service
- Each registered service is assigned a name which clients use to find the service
 - The clients are unaware of the actual, physical location of the service
- Included as part of the Java platform is a RMI registry application called *rmiregistry*, which can be set up to listen for incoming connections

Defining functionality of a RMI service

- The first step in creating a RMI service is defining its functionality in an interface that implements the *Remote* interface
- *Remote* is a *marker* interface that does not have any methods of its own
 - Its sole purpose is to “tag” objects so that they can be identified as RMI services
- This sub-interface of *Remote* determines which methods of the RMI service clients may call
 - A RMI service may have many public methods, but only those declared in a remote interface can be invoked remotely
 - Other public methods may be invoked only from within the virtual machine where the object lives

Implementing a RMI service

- Besides implementing methods of the RMI service interface, the class that serves as a RMI service implementation should extend the *UnicastRemoteObject* class
- This class provides all the core RMI functionality, such as exporting a RMI service and obtaining a client that communicates with the service
- Extending from *UnicastRemoteObject* is the only RMI-specific code that needs to be written for a service implementation
 - Beyond that, there is no actual networking code required
- Once a service implementation exists, the *rmic* tool, which ships with the JDK, is used to create stubs and skeletons
 - Use of *rmic* with RMI is deprecated in newer versions of Java, since stubs and skeletons are generated dynamically

Binding RMI services

- The process of associating a RMI service with a name in a registry is called *binding*
- Binding can be performed through static methods of the *Naming* class:
 - **void bind(String name, Remote obj):** binds the specified name to a remote object
 - **void rebind(String name, Remote obj):** rebinds the specified name to a new remote object; any existing binding for the name is replaced
 - **void unbind(String name):** destroys the binding for the specified name that is associated with a remote object

Binding RMI services (cont')

- **String[] list(String registryName):** returns an array of the names bound in the specified registry
- The name of each RMI service is specified in URL format of the form *rmi://host:port/name*
- The *host* and *port* parameters are optional: the host defaults to the local host, while to port defaults to 1099
 - The *rmi* scheme specification is also optional
- The object that actually binds and instantiates a RMI service is called a *RMI server*

Invoking a RMI service

- To invoke a RMI service, the client application needs only to obtain an object reference to the remote interface
 - It does not need to be concerned with how messages are sent or received, or where the service is located
- To find the service initially, a lookup in the RMI registry is made: the client application invokes the following static method of the *Naming* class
 - **Remote lookup(String name)**: returns a stub for the RMI service associated with the specified name

RMI example – remote interface definition

- Note that each method in the remote interface needs to declare *RemoteException* in its list of possible exceptions

```
class DivisionByZero extends Exception {  
    public DivisionByZero(String reason) {  
        super(reason);  
    }  
}
```

```
public interface Calculator extends Remote {  
    double add(double x, double y) throws RemoteException;  
    double subtract(double x, double y) throws RemoteException;  
    double multiply(double x, double y) throws RemoteException;  
    double divide(double x, double y)  
        throws RemoteException, DivisionByZero;  
}
```


RMI example – RMI service implementation

```
public class CalculatorImpl
    extends UnicastRemoteObject implements Calculator {
    private static final double EPSILON = 0.0001;

    // an implicit constructor that throws
    // RemoteException needs to exist
    protected CalculatorImpl() throws RemoteException { }

    @Override
    public double divide(double x, double y)
        throws RemoteException, DivisionByZero {
        if (Math.abs(y) < EPSILON)
            throw new DivisionByZero("Cannot divide " + x + " by zero");
        return x / y;
    }

    ...
}
```

- The compiled service implementation can be fed to the *rmic* tool to produce the stub and skeleton

RMI example – running the RMI server

- The role of a RMI server is to register a RMI service with the running instance of a RMI registry

```
public class CalculatorServer {  
    public static void main(String[] args) {  
        try {  
            String host = args.length >= 1 ? args[0] : "";  
            String port = args.length >= 2 ? args[1] : "";  
            // build the name  
            String name = String.format("//%s:%s/Calculator", host, port);  
            // bind an instance of the calculator to this name  
            Naming.rebind(name, new CalculatorImpl());  
        } catch (Exception e) {  
            e.printStackTrace();  
        }  
    }  
}
```

RMI example – client implementation

```
public class CalculatorClient {
    public static void main(String[] args) {
        String host = args.length >= 1 ? args[0] : "";
        String port = args.length >= 2 ? args[1] : "";
        try {
            // acquire a reference to the stub
            String name = String.format("//%s:%s/Calculator", host, port);
            Remote remote = Naming.lookup(name);
            CalculatorImpl_Stub calc = (CalculatorImpl_Stub)remote;
            // invoke some operations
            System.out.printf("%f + %f = %f\n", 5.0, 12.0, calc.add(5, 12));
            try {
                System.out.printf("%f / %f = %f\n", 5.0, 0.0, calc.divide(5, 0));
            } catch (DivisionByZero e) {
                System.out.println(e.getMessage());
            }
        } catch (Exception ex) {
            ex.printStackTrace();
        }
    }
}
```

- Console: 5,000000 + 12,000000 = 17,000000
Cannot divide 5.0 by zero

Distributed programming with Java

- Introduction
- RMI
- **Serialization**
- The Reflection API
- Case study – mobile agents

Serialization

- *Serialization* is the process of converting a set of object instances that contain references to each other into a linear stream of bytes
 - It is the mechanism used by RMI to pass objects between JVMs
- The serialization process of an object preserves:
 - The class name and signature of the class
 - The values of all non-static object's fields
 - The closure of any other objects referenced from the initial objects
- The reversed process – creating a set of objects from a stream of data – is called *deserialization*

Using serialization

- The three most common uses of serialization are:
 - *Persistence*: by using *FileOutputStream*, the object stream can automatically be written to a file
 - *A copy mechanism*: by using *ByteArrayOutputStream*, the object stream is written to a byte array in memory, which can then be used to create duplicates of the original object
 - *Communication*: by using a stream that comes from a socket, objects can automatically be sent over the wire to the receiving socket

Serialization requirements

- To allow an object to be serializable, its class should implement the *Serializable* interface
 - *Serializable* is a “marker” interface – it contains no elements, but simply tags the class as serializable
- A *NotSerializableException* is thrown if serialization is tried on non-serializable objects
- Serializability is inherited: it only needs to be implemented once along the class hierarchy
 - Most Java classes are serializable

The *transient* keyword

- In order to serialize an object, all of its fields have to be serializable as well
 - Serialization does not care about access modifiers, such as *private*
- To serialize a class with non-serializable fields, mark the properties with the *transient* keyword
 - The *transient* keyword prevents the data from being serialized
- During deserialization, transient properties are initialized to their default values (0, *null*, etc.)
- The keyword can also be used with serializable fields, e.g. for performance reasons

Writing and reading objects

- *ObjectOutputStream* writes primitive data types and graphs of Java objects to an output stream
- It offers the following method for writing an object to a stream:
 - **void writeObject(Object obj)**
- *ObjectInputStream* deserializes primitive data and objects previously written using an *ObjectOutputStream*
- It offers the following method for reading an object from a stream:
 - **Object readObject()**

Persistent list of numbers – example

```
public class NumList implements Serializable {
    private List<Integer> numbers;
    private int maxNum;
    private transient int minNum;

    public NumList() {
        numbers = new ArrayList<Integer>();
        maxNum = Integer.MIN_VALUE;
        minNum = Integer.MAX_VALUE;
    }

    public void add(int num) {
        numbers.add(num);
        // remember max and min values
        if (maxNum < num)
            maxNum = num;
        if (minNum > num)
            minNum = num;
    }

    @Override public String toString() {
        return String.format("%s (min:%d; max:%d)",
            numbers, minNum, maxNum);
    }
}
```

Persistent list of numbers – example (cont')

```
public class SerializationTest {  
  
    private static NumList load(String fileName) {  
        ObjectInputStream in = null;  
        try {  
  
            File file = new File(fileName);  
            if (!file.exists())  
                return new NumList();  
            in = new ObjectInputStream(new FileInputStream(file));  
            return (NumList)in.readObject();  
  
        } catch (Exception e) {  
            return new NumList();  
        } finally {  
            if (in != null)  
                try {  
                    in.close();  
                } catch (IOException e) { }  
        }  
    }  
}
```

Persistent list of numbers – example (cont')

```
private static void save(NumList nums, String fileName) throws IOException {
    ObjectOutputStream out = new ObjectOutputStream(
        new FileOutputStream(fileName));
    try {
        out.writeObject(nums);
    } finally {
        out.close();
    }
}

public static void main(String[] args) throws IOException {
    NumList nums = load("numbers.dat");
    System.out.println("Existing data: " + nums);

    for (int i = 0; i < 3; i++)
        nums.add((int) (Math.random() * 9) + 1);

    System.out.println("After adding: " + nums);
    save(nums, "numbers.dat");
}
```

- First run: Existing data: [] (min:2147483647; max:-2147483648)
After adding: [7, 5, 2] (min:2; max:7)
- Second run: Existing data: [7, 5, 2] (min:0; max:7)
After adding: [7, 5, 2, 9, 4, 2] (min:0; max:9)

Version control

- All serializable classes are automatically given a version identifier
- This identifier is saved along with the object and automatically updated whenever the object's class changes
 - E.g. when a new field is added
- Version identifiers are compared during deserialization: if the version of the class does not equal the version of the object in the stream, an exception is thrown
- To control the versioning system, developers simply need to provide the static *serialVersionUID* field manually and ensure it is always the same, unless such changes are made to the class which invalidate previously serialized objects

Distributed programming with Java

- Introduction
- RMI
- Serialization
- The Reflection API
- Case study – mobile agents

The Reflection API

- *Reflection* is the process by which software can observe and modify (its own) program structure and behavior at runtime
- It allows inspection of classes and their elements, as well as instantiation of new objects and invocations of methods at runtime without knowing the actual names at compile time
- In Java, the access to this *object metadata* is available through an immutable instance of *java.lang.Class*
- The Java Reflection API, available under the *java.lang.reflect* package, is most commonly used by:
 - Serialization and RMI
 - Class browsers and visual development environments
 - Debuggers and test tools that, for example, need to access private properties of a class

Retrieving the object metadata

- There are several ways of retrieving the *Class* instance:
 - If an instance of an object is available, the simplest way to get its *Class* is to invoke its *getClass()* method
 - If the type is available but there is no instance, then it is possible to obtain a *Class* by appending “.class” to the name of the type; this approach also works for primitive data types
 - If the fully-qualified name of a class is available, it is possible to get the corresponding *Class* using the static method *Class.forName()*
- For primitive types, e.g. *int*, use *int.class*, or *Integer.TYPE*

Examining class modifiers and types

- The following example demonstrates how to examine class modifiers, generic type parameters, inheritance path, and annotations

```
public class ClassDeclarationSpy {
    public static void main(String[] args) {
        try {
            Scanner scanner = new Scanner(System.in);
            out.print("Class name? ");
            String name = scanner.nextLine();

            Class<?> c = Class.forName(name);
            out.println("Modifiers:\n\t" +
                Modifier.toString(c.getModifiers()));

            out.println("Type Parameters:");
            TypeVariable<?>[] tv = c.getTypeParameters();
            if (tv.length == 0)
                out.println("\tNo Type Parameters");
            else
                for (TypeVariable<?> t : tv)
                    out.println("\t" + t.getName());
        }
    }
}
```

Examining class modifiers and types (cont')

```
out.println("Implemented interfaces:");
Type[] intfs = c.getGenericInterfaces();
if (intfs.length == 0)
    out.println("\tNo Implemented Interfaces");
else
    for (Type intf : intfs)
        out.println("\t" + intf);

out.println("Inheritance path:");
printAncestors(c);

out.println("Annotations:");
Annotation[] ann = c.getAnnotations();
if (ann.length == 0)
    out.println("\tNo Annotations");
else
    for (Annotation a : ann)
        out.print("\t" + a);

} catch (Exception e) {
    e.printStackTrace();
}
}
```

Examining class modifiers and types (cont')

```
private static void printAncestors(Class<?> c) {  
    Class<?> ancestor = c.getSuperclass();  
    if (ancestor != null) {  
        out.println("\t" + ancestor.getCanonicalName());  
        printAncestors(ancestor);  
    }  
}
```

■ Console:

Class name? java.lang.String;

Modifiers:

public abstract final

Type Parameters:

No Type Parameters

Implemented interfaces:

interface java.lang.Cloneable

interface java.io.Serializable

Inheritance path:

java.lang.Object

Annotations:

No Annotations

Discovering class members

- There are two categories of methods provided in *Class* for accessing fields, methods, and constructors:
 - Methods which enumerate all members
 - Methods which search for a particular member
- Additionally, there are distinct methods for accessing members declared directly on the class versus methods which search the super-interfaces and super-classes for inherited members
- Finally, some methods of *Class* can look for public members only, while other can also access private and protected members

Discovering class members – example

- The following example demonstrates how to list all members of a class, declared directly on the class

```
package test;

public class SampleClass {
    private class InnerClass {
        private int k;
        private InnerClass(int k) { this.k = k; }
    }

    private int privateField;
    public InnerClass publicField;

    protected SampleClass() { }

    public SampleClass(int n) {
        privateField = n;
        publicField = new InnerClass(n);
    }

    public int getValue() {
        return privateField + publicField.k;
    }
}
```

Discovering class members – example (cont')

```
public class ClassMemberSpy {
    private static void printClass(Class<?> c) {
        out.println("Class: " + c.getName());
        out.println("Package: " + c.getPackage());

        printMembers("Constructors:", c.getDeclaredConstructors());
        printMembers("Methods:", c.getDeclaredMethods());
        printMembers("Fields:", c.getDeclaredFields());

        out.println("Inner classes:");
        Class<?>[] inner = c.getDeclaredClasses();
        if (inner.length == 0)
            out.println("\tNone");
        else
            for (Class<?> cls : inner)
                printClass(cls);
    }

    private static void printMembers(String msg, Member[] members) {
        out.println(msg);
        if (members.length == 0)
            out.println("\tNone");
        else
            for (Member m : members)
                out.println("\t" + m);
    }
}
```

Discovering class members – example (cont')

```
public static void main(String[] args) throws ClassNotFoundException {  
    Scanner scanner = new Scanner(System.in);  
    out.print("Class name? ");  
    String name = scanner.nextLine();  
    Class<?> c = Class.forName(name);  
    printClass(c);  
}  
}
```

■ Console:

Class name? test.SampleClass

Class: test.SampleClass

Package: package test

Constructors:

protected test.SampleClass()

public test.SampleClass(int)

Methods:

public int test.SampleClass.getValue()

Fields:

private int test.SampleClass.privateField

public test.SampleClass\$InnerClass test.SampleClass.publicField

Inner classes:

Class: test.SampleClass\$InnerClass

Package: package test

Constructors:

...

Distributed programming with Java

- Introduction
- RMI
- Serialization
- The Reflection API
- Case study – mobile agents

Case study – mobile agents

- *Agent technology* represents one of the most consistent approaches to distributed software development
- It employs a distributed network of autonomous, executable software entities called *agents*
- An agent is considered to be *mobile* if it can move from one computer in a network to another
- A common use of mobile agents is in processing of large data sets:
 - Retrieving a large data set from a remote database to a computer hosting the processing software can be too expensive, or even impossible
 - A much better approach is to send a (usually small) agent to the target computer and perform the processing on-the-spot

Agent server

- Each computer that wants to accept mobile agents hosts an *agent server*
- The server accepts a serialized form of a mobile agent, and then deserializes and runs it
- To make the server as general as possible, the Reflection API is used to run the agent
 - The server will look for and invoke the following method: *void onArrival(ServerSocket host)*

Agent server implementation

```
public class AgentServer extends Thread {
    public static final int PORT = 6060;

    @Override public void run() {
        ServerSocket socket = null;
        try {
            socket = new ServerSocket(PORT);
            while (true) {
                Socket client = socket.accept();
                new AgentHandler(socket, client).start();
            }
        } catch (Exception ex)
        {
            ex.printStackTrace();
        } finally
        {
            if (socket != null)
                try
                {
                    socket.close();
                } catch (IOException e) { }
        }
    }
}
```

Agent server implementation (cont')

```
public class AgentHandler extends Thread {
    private ServerSocket server;
    private Socket client;

    public AgentHandler(ServerSocket server, Socket client) { ... }

    @Override public void run() {
        ObjectInputStream in = null;
        try {
            in = new ObjectInputStream(client.getInputStream());
            // deserialize the agent and get its class
            Object agent = in.readObject();
            Class<?> c = agent.getClass();
            // look for the "onArrival" method
            // with one parameter of type ServerSocket
            Method m = c.getMethod("onArrival", ServerSocket.class);
            // invoke the method
            m.invoke(agent, server);
        } catch (Exception e) {
            e.printStackTrace();
        } finally {
            if (in != null)
                try {
                    in.close();
                } catch (IOException e) { }
        }
    }
}
```

The mobile agent

- The source code of the mobile agent is very simple:

```
public class MobileAgent implements Serializable {  
    public void moveTo(String host, int port) {  
        Socket socket = null;  
        ObjectOutputStream out = null;  
        try {  
            socket = new Socket(host, port);  
            out = new ObjectOutputStream(socket.getOutputStream());  
            out.writeObject(this);  
        } catch (Exception e) {  
            e.printStackTrace();  
        } finally {  
            // close socket and output stream  
        }  
    }  
  
    public void onArrival(ServerSocket host) {  
        // perform data processing on the current host  
    }  
}
```